

Yes — **overall your setup is right**, and more importantly, it is now a **coherent architecture**.

What is right about it

You have the correct building blocks for this kind of product:

- **IIS + domain + SSL** for public access
- **Dockerized app** for deployment consistency
- **Postgres** for transactional data
- **Qdrant** for vector search
- **Ollama** for local AI inference
- **WSL bridge/relay** to connect Dockerized Linux services to Windows Ollama
- **Cloudflare + domain routing** in front of the app

That is a valid setup for an AI marketplace MVP.

What is good in your design

The strongest parts of your setup are:

- you are **not depending on OpenAI**
- you are keeping **embeddings local**
- you are using **Qdrant for search**
- you already have **Arabic as a priority**
- you have a **production-accessible domain**
- your app, DB, and vector DB are already separated properly

Architecturally, that is the right direction.

Where your setup is weak

The main weakness is not the application design.

3/14/26 5:44:43 AM

It is the **operational complexity caused by Windows + WSL + Docker + Ollama together.**

Right now you have this chain:

- IIS on Windows
- App in Docker/WSL
- Ollama on Windows
- relay in WSL
- container calling host alias through relay

This works, but it is fragile.

The weak points are:

- Ollama binding may change or restart incorrectly
- WSL relay may stop after reboot/session issues
- host networking between Docker and Windows is awkward
- troubleshooting becomes harder than it should be

So the setup is **correct**, but **more complicated than ideal**.

My honest assessment

For now

Yes, your setup is good enough to launch and test users.

For long-term stability

The cleanest architecture would be one of these:

- **Option A:** move everything to **Linux**
- **Option B:** keep Windows, but run Ollama and app in the same Linux/Docker environment
- **Option C:** separate AI inference onto another dedicated box later

What I think is the “right” setup for your current phase

For your current server and business stage, the right setup is:

- **Monetchat app** in Docker
- **Postgres** in Docker
- **Qdrant** in Docker
- **Ollama with qwen2.5:7b**
- **nomic-embed-text** for embeddings
- **llava** only when needed
- **IIS reverse proxy with SSL on chatsoq.com**

That is a sensible MVP stack.

My verdict

Yes, the whole setup is conceptually right.

But the Windows/WSL/Ollama bridge is the one part I would eventually simplify.

So I'd describe it like this:

- **Product architecture:** right
- **Technology choices:** right
- **Deployment approach for MVP:** acceptable
- **Infrastructure elegance:** not ideal, but workable
- **Best long-term improvement:** move to a cleaner Linux-based deployment

