



For **your project (Ollama + AI search + embeddings + Qdrant + chat)**, the decision should be based on **cost, flexibility, and reliability**.

Here is the **honest recommendation**:

1 ☐ **Best Overall Choice:**

RunPod

Go with RunPod.

Reasons:

- Built specifically for **AI workloads**
- Very easy to deploy **GPU pods**
- Pay **per second**
- Can scale GPUs up/down anytime
- Excellent for **Ollama hosting**
- Used by many AI startups

Typical cost:

- **RTX 4090**: ~\$0.50–0.60/hr
- Monthly if always on: **\$350–\$450**

Website:

<https://runpod.io>

2 ☐ **Cheapest Option:**

Vast.ai



If your priority is **lowest cost**.

Pros:

- Sometimes **\$0.25–0.40/hr**
- Huge GPU marketplace

Cons:

- Less stable
- Some hosts unreliable
- Setup more manual

Website:

<https://vast.ai>

3 ☐ **Not Recommended for AI GPUs:**

DatabaseMart

Reasons:

- Very expensive for GPU
- Hardware usually older
- Designed more for **traditional servers**

You could still keep your **existing server there for the application, but not the GPU.**

Best Architecture for Your System

Server 1 (DatabaseMart / normal VPS)



Runs:

- Next.js app
- API
- Qdrant
- Database

Server 2 (RunPod GPU)

Runs:

- Ollama
- LLM models
- embeddings

Your app calls the GPU server like:

```
Next.js App
  ↓
API Server
  ↓
Ollama GPU Server
```

My Final Recommendation

Start with:

RunPod RTX 4090

Cost:

~ \$350–\$450/month

Power:

Enough for:

- Llama 3
- Mixtral



- DeepSeek
- Qwen
- embeddings
- Ollama

Important Advice for Your Project

Do **NOT** rent:

- H100
- H200
- B200

These are **\$2000–\$4000/month GPUs** and unnecessary for your stage.
