

sell sell marketplace updates:

SHAZ EJ



SellSell / PickPic RunPod Recovery & Deployment Guide

Purpose

This document explains how to recreate a RunPod GPU pod after accidental deletion while preserving Ollama models stored in the Network Volume.

Existing Resources

Network Volume

Volume Name:

ollama-models

Volume Size:

200 GB

Data Center:

US-MO-2

Important:

Never delete this volume. The pod can be recreated at any time, but the volume contains the Ollama models and other persistent data.



Step 1: Create a New Pod

Login to RunPod:

<https://console.runpod.io>

Navigate to:

Pods → Deploy Pod → GPU Pod

Step 2: Select Region

Choose:

US-MO-2

This matches the location of the existing Network Volume.

If no GPU is available in US-MO-2:

1. Temporarily select Any Region.
2. Deploy the pod.
3. Ensure the Network Volume can still be attached.

Step 3: Select GPU

Recommended GPU:

RTX 4090

Alternative options:

- RTX 5090
- L40S
- RTX PRO 4500

Avoid expensive enterprise GPUs unless required:

- H100
- H200
- B200

Current preferred configuration:

6/3/26 2:19:05 PM



RTX 4090
\$0.69/hr

Step 4: Configure Deployment

Pod Name

Set:

sellsell-ollama

Pod Template

Use:

RunPod Pytorch 2.4.0

Template Example:

runpod/pytorch:2.4.0-py3.11-cuda12.4.1-devel-ubuntu22.04

Pricing

Select:

On-Demand

Enable

✓ SSH Terminal Access

Disable

X Start Jupyter Notebook

(Not required)

6/3/26 2:19:05 PM



Leave Disabled

X Encrypt Volume

Step 5: Attach Existing Network Volume

Locate the Network Volume section.

Select:

ollama-models (200 GB)

DO NOT create a new volume.

DO NOT leave volume attachment empty.

This is the most important step.

Step 6: Container Disk

Recommended:

30 GB

or

50 GB

No large container disk is needed because models are stored on the Network Volume.

Step 7: Exposed Ports

Add:

11434

Purpose:



Ollama API

Optional:

8000

Purpose:

Transcription Service

Step 8: Environment Variables

Add:

OLLAMA_HOST=0.0.0.0:11434

Step 9: Deploy

Click:

Deploy

Wait until the pod status becomes:

Running

Step 10: Connect via SSH

Copy the SSH command from the pod page.

Example:

ssh root@YOUR_POD_IP -p YOUR_PORT

Connect using Terminal on macOS.

Step 11: Verify Network Volume

Inside the pod run:

6/3/26 2:19:05 PM



df -h

Then:

ls -lah /workspace

Then:

du -sh /workspace/*

Verify the existing files are visible.

Step 12: Verify Ollama

Run:

ollama list

Expected output:

qwen2.5:14b-instruct

If the model appears, the recovery is successful.

Step 13: Start Ollama

Run:

ollama serve

Leave this process running.

Step 14: If Ollama Is Missing

Install Ollama:

curl -fsSL <https://ollama.com/install.sh> | sh

Then start:



ollama serve

Check:

ollama list

Step 15: Create Local SSH Tunnel

On Mac:

```
ssh -N -L 11435:localhost:11434 root@YOUR_POD_IP -p YOUR_PORT
```

Keep this terminal open.

Step 16: Test Ollama

Open another terminal:

```
curl http://localhost:11435/api/tags
```

Expected result:

JSON containing installed Ollama models.

Step 17: Run SellSell Locally

Navigate to project:

```
cd /Users/shazej/Documents/askar/1705/pickpick-slsI
```

Start Docker:

```
docker compose up
```

Step 18: Environment Variables

Verify:

```
OLLAMA_BASE_URL=http://host.docker.internal:11435
```

6/3/26 2:19:05 PM



AI_BASE_URL=http://host.docker.internal:11435

OLLAMA_CHAT_MODEL=qwen2.5:14b-instruct

Step 19: Final Validation

Test:

1. Chat functionality
2. AI-generated descriptions
3. Product search
4. Seller listing flow
5. Ollama API response

If all tests pass, the system is fully restored.

Important Reminder

Pods are temporary.

Network Volumes are permanent.

Always protect:

ollama-models (200 GB)

This volume contains the persistent AI models and should never be deleted.