

sell sell marketplace updates:

SHAZ EJ



SellSell / PickPic Production AI Infrastructure Update

Date: 03 June 2026

Objective

Restore AI services after RunPod GPU pod migration and reconnect production to RunPod-hosted Ollama models.

Problem Identified

The previous RunPod GPU pod became unavailable and a replacement pod had to be created.

Issues encountered:

- 1. RTX 4090 was unavailable on RunPod.**
- 2. New pod was created using a different GPU.**
- 3. Ollama service was not running on the new pod.**
- 4. HTTP port 11434 was not exposed initially.**
- 5. Production application could not communicate with the Ollama endpoint.**
- 6. AI-powered features became unavailable.**

Actions Completed

1. RunPod Pod Recovery



Verified persistent network storage was mounted successfully.

Existing data preserved:

- Ollama models
- Custom model configurations
- Embedding models
- Vision models
- Transcription service files

Storage verification showed:

- Models directory preserved
- Ollama directory preserved
- Transcription service preserved

2. Ollama Restoration

Installed missing dependency:

bash apt-get install -y zstd

Executed:

bash bash /workspace/start-ollama.sh

Verified Ollama service:



bash ps aux | grep ollama

Verified API:

bash curl http://localhost:11434/api/tags

3. Model Verification

Confirmed available models:

- qwen2.5:14b-instruct**
- qwen2.5:14b-instruct-q8_0**
- qwen-coder**
- deepseek-coder**
- nomic-embed-text**
- llava**

These are the core models used by SellSell.

4. Public API Exposure

Updated RunPod configuration:

HTTP Ports:

text 8888 11434



Result:

Public Ollama endpoint became available:

text <https://xww6sf0m82ow0w-11434.proxy.runpod.net>

5. Production Reconfiguration

Updated production environment:

```
env AI_BASE_URL=https://xww6sf0m82ow0w-11434.proxy.runpod.net
OLLAMA_BASE_URL=https://xww6sf0m82ow0w-11434.proxy.runpod.net
OLLAMA_CHAT_MODEL=qwen2.5:14b-instruct-q8_0
```

Restarted production:

```
bash docker compose up -d
```

6. Connectivity Validation

Verified from production container:



```
bash docker exec Sellsell_app curl  
https://xww6sf0m82ow0w-  
11434.proxy.runpod.net/api/tags
```

Result:

- Success
- All models visible
- Production can communicate with RunPod

7. Application Validation

Verified:

- Next.js application starts successfully
- PostgreSQL connected
- Redis connected
- Qdrant connected
- Worker service started
- AI endpoint reachable

Application log:

text ✓ Ready Redis connected successfully

Current Production Architecture



Production Server

DigitalOcean

Services:

- Next.js Application
- PostgreSQL
- Redis
- Qdrant
- Background Workers

RunPod GPU Server

Services:

- Ollama
- Qwen 14B Chat Model
- Qwen Coder
- Nomic Embeddings
- Llava Vision

Communication

Production Server

↓ HTTPS

RunPod Ollama API

text <https://xww6sf0m82ow0w-11434.proxy.runpod.net>



Remaining Issue Identified

Worker logs show an analytics-related error:

text Argument productCode is missing

Affected component:

- Analytics Worker

Impact:

- Does NOT affect search**
- Does NOT affect chat**
- Does NOT affect listing creation**
- Does NOT affect moderation**
- Does NOT affect embeddings**

Impact limited to analytics logging jobs.

Requires separate code fix.

Current System Status

Operational

- AI Chat**
- Ollama**



- Qwen 14B
- Embeddings
- Vision Model
- Product Search
- Semantic Search
- Moderation Pipeline
- Qdrant
- PostgreSQL
- Redis
- Background Workers

Needs Follow-up

- Analytics Worker productCode bug
- Qdrant client version warning
- Production smoke test of AI listing generation
- Production smoke test of conversational search

Recommended Next Steps

1. Test conversational search in production.
2. Test seller listing flow.
3. Test AI-generated descriptions.
4. Test image analysis flow.
5. Fix analytics worker productCode bug.
6. Fix Qdrant client/server version mismatch warning.
7. Run full browser QA pass.
8. Prepare production launch checklist.



Overall Status

Production AI infrastructure has been successfully restored.

RunPod Ollama is operational.

Production application is connected to RunPod successfully.

Core marketplace functionality is available for testing.