




MONETCHAT — COMPLETE TECHNICAL AUDIT

CTO-Level Review | March 2026

SYSTEM ARCHITECTURE DIAGRAM



AI_PROVIDER=ollama → Local Ollama (fallback)
AI_PROVIDER=openai → OpenAI (production)



1. SYSTEM OVERVIEW — WHAT IT IS

Monetchat is a **conversational commerce marketplace** for Kuwait (with Saudi Arabia planned). Users buy and sell by chatting with an AI — no browse/filter UI. The AI interprets text, voice (Arabic/English), and images to either find products or list them.

Core Value Proposition:

- Buyer types/speaks "أبي سيارة عائلية بسعر مناسب" → AI expands the query, embeds it, vector-searches Qdrant, and shows matching listings with seller contact buttons
- Seller uploads a photo of their item → AI analyzes it via GPT-4o Vision, generates title/description/category in both Arabic and English, creates a draft listing for review

2. CORE ARCHITECTURE BREAKDOWN

Backend: Next.js 15 API Routes

All backend logic lives in `src/app/api/`. There is no separate Express/Fastify server — Next.js handles everything. The `server.js` file exists but is unused by the main app (only referenced for the worker). This is a **clean, modern monorepo** pattern.

Route Count: 18 API endpoints across auth, chat, products, upload, geo, AI, admin, seller, health.

Frontend: Next.js App Router + React 18

Full-page chat interface is the homepage (`(marketplace)/page.tsx`). The entire UX is built around `chat-interface.tsx` — a 1000+ line React component managing SSE streaming, voice recording, image upload, product display, and draft listings all in one place.

State management: React Context (`app-mode-context.tsx`, `language-context.tsx`) + local `useState/useRef` — no Redux, no Zustand. Appropriate for current scale.

Database: PostgreSQL via Prisma

20 models (not 15 as documented — schema shows `RefreshToken`, `Notification`, `NotificationJob`, `Attachment` were added):

Model	Purpose
User	Auth, preferences, locale
RefreshToken	Refresh token rotation (revocable)
Seller	Business profile, phone, WhatsApp
Product	Listings with dual-language fields



Model	Purpose
ProductImage	S3 URLs per product
Category	Hierarchical (parent/child)
Country	KW/SA with currency/phone codes
Region	Sub-national areas
ChatSession	Conversation threads
ChatMessage	Messages with product ID arrays
ProhibitedKeyword	Content moderation list
Role + Permission + RolePermission	RBAC
SearchLog	Analytics (query, type, latency)
ContactLog	Call/WhatsApp tracking
AuditLog	All entity changes with IP
Notification	User notifications
Attachment	Generic file registry
NotificationJob	Queue job tracking

Indexes exist on: `products(sellerId, categoryId, countryCode, status, price, createdAt), users(email, countryCode), chatSessions(userId, lastMessageAt), auditLogs(entityName+entityId, userId).`

Vector DB: Qdrant with Hybrid Search

This is the most sophisticated piece of the system. The Qdrant collection uses:

Dense vector: size=192 (Ollama nomic-embed-text dimensions) ← CRITICAL BUG
Sparse vector: BM25 with IDF modifier
Fusion: RRF (Reciprocal Rank Fusion)

⚠ CRITICAL BUG: `VECTOR_SIZE = 192` is hardcoded in `qdrant/client.ts` for Ollama's `nomic-embed-text`. But `ai-service.ts` uses `AiFactory.getProvider().createEmbedding()` — if `AI_PROVIDER=openai`, OpenAI's `text-embedding-3-small` produces **1536 dimensions**, not 192. This will cause a dimension mismatch error when indexing or searching with OpenAI embeddings. The Qdrant collection would need to be recreated with `size: 1536` when running with OpenAI.

AI/LLM Integration

Controlled by `AI_PROVIDER` env var via the **Factory pattern**:

- `AI_PROVIDER=ollama` → `OllamaProvider` (local, free, slower, CPU)
- `AI_PROVIDER=openai` → `OpenAiProvider` (cloud, paid, faster, GPU)

Providers expose a common interface:

- `createChatCompletion()` — standard chat
- `createChatCompletionStream()` — streaming with tool calls




- `createEmbedding()` — text → vector
- `getDefaultChatModel()` — model name

Voice transcription is always Groq (whisper-large-v3) regardless of `AI_PROVIDER` — Ollama doesn't support audio natively.

3. AI & SEARCH SYSTEM — DEEP ANALYSIS

The Full Query Pipeline (Step by Step)

User Message → `/api/chat` (POST)





How Well Does It Handle Imprecise Queries?

Strengths:

- **Bilingual expansion:** "family car" → Arabic+English keywords via GPT-4o preprocessing — this is excellent
- **Hybrid search:** Combines semantic (meaning) + BM25 (exact keywords) — covers both fuzzy and exact matches
- **Metadata reranking:** Brand/model exact matches get boosted — reduces irrelevant results
- **Intent detection:** Distinguishes buy/sell/clarify autonomously

Weaknesses & Gaps:

Gap	Severity	Details
No score threshold	Medium	Returns results even with very low confidence — user sees irrelevant items
Reranker is heuristic-only	Medium	No cross-encoder reranking (Cohere, BGE) — purely rule-based boosts
<code>preprocessSearchQuery</code> is a blocking LLM call	High	Every search makes 2 GPT-4o calls (preprocess + main). At scale this doubles cost and latency
No embedding cache	Medium	Same query re-embeds every time — no Redis cache for embeddings
Arabic BM25 needs stemming	Medium	"سيارات" (cars) won't match "سيارة" (car) — Arabic morphology is complex
<code>Qdrant indexing_threshold: 10000</code>	Low	HNSW index only builds after 10k points — sub-10k searches use brute force (fine for MVP)
<code>query: { fusion: 'rrf' }</code> uses default weights	Low	Dense and BM25 have equal weight — likely want to tune dense heavier for semantic queries

4. CODE QUALITY & COMPLETENESS

Fully Implemented

- Auth system (register/login/logout/JWT/refresh tokens/RBAC)
- Full chat pipeline with SSE streaming and tool calling
- Product CRUD with Qdrant indexing
- Seller flow (upload → GPT-4o vision → draft → publish)
- Voice transcription via Groq Whisper
- Image-to-search and image-to-listing flows
- Admin analytics routes
- Content moderation (prohibited keyword check)
- Health check for all 4 services



- Retry logic with exponential backoff on rate limits
- Hybrid search with BM25 + dense vectors + RRF

Partially Implemented ⚠

Item	Status	Risk
Localization	60% — basic Arabic strings done, error messages missing	Medium
Mobile responsiveness	CSS RTL support done, layout not fully tested	Medium
Dark mode	Not implemented	Low
Location picker UI	Missing component	Low
E2E tests	Spec files exist but not populated	High for production
<code>executeCreateListing()</code> in <code>route.ts</code>	Defined but never called — the route emits <code>listing_draft</code> SSE, actual product creation is done from frontend <code>POST /api/products</code>	Medium — code dead weight
Notification system	Model + jobs table exist, no delivery code	Low
Attachment model	Schema exists, no usage	Low

Bugs Found 🐛

Bug 1 — CRITICAL: Qdrant vector dimension mismatch

```
qdrant/client.ts: VECTOR_SIZE = 192 ← Ollama nomic-embed-text
ai-service.ts: AiFactory.getProvider().createEmbedding() ← Could be OpenAI (1536-dim)
```

If `AI_PROVIDER=openai`, `indexProduct()` will throw a Qdrant error because it tries to store a 1536-dim vector in a 192-dim collection. Collection must be created with the correct size per provider.

Bug 2 — MEDIUM: `executeCreateListing` is dead code `route.ts` defines `executeCreateListing()` at line 719 but never calls it. The `create_listing` tool handler at line 372 emits a `listing_draft` SSE event instead of calling this function. The actual listing is created by the frontend via `POST /api/products`. This is intentional by design (user reviews draft first) but the dead function adds confusion.

Bug 3 — MEDIUM: Middleware auth mismatch `src/middleware.ts` protects `/api/products/create` and `/api/products/edit` — but the actual routes are `POST /api/products` and `PUT /api/products/[id]`. The protected path strings don't match the real routes. The routes rely on `getCurrentUser()` inside the handler instead. This means middleware RBAC is ineffective for products.

Bug 4 — MEDIUM: CORS wildcard in middleware

```
'Access-Control-Allow-Origin': '*' // middleware.ts line 68
```



Wildcard CORS with credentials is a security issue. Should be locked to the production domain.

Bug 5 — LOW: s3Key derivation is fragile

```
s3Key: url.split('/').slice(-2).join('/') // route.ts line 781
```

This assumes a specific URL structure. If S3 CDN URL format changes, keys will be wrong and files become orphaned.

Bug 6 — LOW: Chat history limited to 20 messages without pruning strategy

```
take: 20 // route.ts line 103
```

No token counting — if messages are long (image URLs, rich descriptions), the 20-message window could exceed GPT-4o's context limit and cause truncation silently.

Scalability Issues

Issue	Impact
Preprocessing LLM call on every search	Doubles API cost + 500-800ms latency overhead
No connection pooling config for Prisma	Default 5 connections — will hit limits at ~50 concurrent users
Qdrant <code>on_disk_payload: true</code>	Good for memory, slight latency on payload retrieval
Redis used for rate limiting but no session caching	Missed optimization opportunity
S3 in eu-north-1 (Stockholm)	High latency from Kuwait (~200-300ms) — consider me-south-1 (Bahrain)

5. DEPLOYMENT & INFRASTRUCTURE ANALYSIS

Current Setup

```
Windows Server
├── NGINX (SSL + reverse proxy)
├── PM2 (Next.js process manager)
├── Docker Desktop / Docker Engine
│   ├── Postgres:5432
│   ├── Qdrant:6333,6334
│   └── Redis:6379
├── Ollama (Windows daemon, local AI)
└── Next.js app (port 3000 or 9002 dev)
```

Docker Compose — Issues Found



Issue	Risk
Postgres password <code>Monetchatpassword</code> hardcoded in compose	High — should use Docker secrets or env var
No healthcheck for postgres/qdrant/redis services	Medium — app may start before DB is ready
No resource limits (CPU/memory)	Medium — a single query storm can OOM the host
<code>qdrant:latest</code> tag	Low — unpinned version; breaking changes between releases
Worker uses <code>Dockerfile.worker</code> — file not found in repo	High — worker container won't build
App port mapped to 3000 but dev runs on 9002	Low — confusing, needs documentation

Missing Production Readiness

Missing Item	Priority
<code>Dockerfile.worker</code>	High — worker won't deploy
Health checks in docker-compose	High
<code>.env</code> validation on startup (e.g. via <code>zod</code>)	High
Database backup strategy	High
Log aggregation (Pino → file or ELK)	Medium
NGINX config file not in repo	Medium
SSL certificate renewal automation	Medium
Qdrant backup/snapshot config	Medium
Monitoring/alerting (uptime, error rates)	Medium

6. SCALABILITY ANALYSIS

Can It Scale to 10,000 Products? YES

At 10k products Qdrant performs excellently — the `indexing_threshold: 10000` means HNSW index kicks in exactly at this point. PostgreSQL handles 10k rows trivially. S3 is unlimited. No bottleneck.

Estimated search latency: ~200-400ms (embedding + Qdrant + Postgres fetch + rerank)

Can It Scale to 100,000+ Products? WITH CHANGES

Bottleneck	Current State	Fix Needed
Qdrant HNSW	Fine up to millions of vectors	Tune <code>m</code> and <code>ef_construct</code> for accuracy/speed tradeoff
PostgreSQL	No connection pool config	Add <code>PgBouncer</code> or set <code>connection_limit</code> in <code>DATABASE_URL</code>



Bottleneck	Current State	Fix Needed
Preprocessing LLM call	1 GPT-4o call per search	Cache preprocessed queries in Redis (TTL 1h)
Single Qdrant instance S3 in Stockholm	No replication 200ms from Kuwait	Add Qdrant cluster (3 nodes) for HA Move to AWS Bahrain (me-south-1)
Next.js single process (PM2)	CPU-bound at scale	Add PM2 cluster mode (<code>instances: "max"</code>)
No CDN for static assets	Direct S3 reads	Add CloudFront distribution

7. SECURITY ANALYSIS

Risks Found

Risk	Severity	Location	Fix
CORS wildcard *	High	<code>middleware.ts:68</code>	Lock to production domain
Middleware doesn't protect actual product routes	High	<code>middleware.ts:7-11</code>	Fix path patterns to match real routes
Postgres credentials in docker-compose	High	<code>docker-compose.yml:52</code>	Use Docker secrets
<code>nationalId</code> stored in plain text	Medium	<code>schema.prisma:99</code>	Encrypt at application level
No file type validation on upload	Medium	<code>/api/upload/route.ts</code>	Validate MIME type + extension server-side
No max file size enforced server-side	Medium	<code>/api/upload/presign</code>	Add size limit to presigned URL policy
JWT hardcoded 24h expiry, no rotation	Medium	<code>jwt.ts:14</code>	Refresh token rotation is implemented — ensure client uses it
Seller phone numbers in product search results	Low	<code>route.ts:700-702</code>	Confirm intentional (contact is the CTA)
No CSRF protection	Low	All API routes	Add CSRF token for state-changing operations
IP-based rate limiting bypassed by proxy	Low	<code>middleware.ts:22</code>	Trust only specific proxy headers
<code>req.json()</code> with no body size limit	Low	All API routes	Add <code>bodyParser: { sizeLimit: '1mb' } config</code>

What's Done Well Securely

- `httpOnly: true` cookies for JWT — XSS can't steal tokens
- `secure: true` in production
- Passwords hashed (not stored plain — `passwordHash` field)



- Role-based access control with middleware
- Rate limiting on auth routes (stricter) and all API routes
- Prisma ORM — no raw SQL injection risk
- Content moderation for prohibited keywords
- Soft delete (`deletedAt`) — data not permanently lost
- Audit logs with IP address tracking

8. BUSINESS LOGIC UNDERSTANDING

User Journeys

Buyer:

1. Open site → full-screen chat
2. Type/speak/upload → AI searches → product cards appear inline
3. "See all N results" → overlay with full list
4. Click product → detail dialog with Call + WhatsApp buttons
5. Contact seller directly (no escrow, no payment, P2P)

Seller:

1. Upload photo in chat
2. Intent picker: "Sell this item"
3. AI analyzes: title, description, category (bilingual)
4. Draft card appears — review, edit, set price
5. Publish → DB + Qdrant indexed

Revenue Model (Inferred)

The V1 code has **no payment integration, no subscription gates, no listing fees**. Revenue model is unclear from code — likely planned as:

- Freemium listing (free to list, pay for featured)
- Commission (future)
- B2B API access (future)

Target Users

- Primary: Arabic-speaking buyers/sellers in Kuwait
- Secondary: Bilingual English/Arabic users
- Tertiary: Saudi Arabia market (SA country + SAR currency seeded)

9. IMPROVEMENT ROADMAP



Priority 1 — CRITICAL (Do Before Any Production Traffic)

1a. Fix Qdrant vector dimension `VECTOR_SIZE` must match the embedding provider.
Implement:

```
AI_PROVIDER=openai → VECTOR_SIZE=1536  
AI_PROVIDER=ollama → VECTOR_SIZE=192
```

Read from env or provider interface at collection creation time.

1b. Fix middleware route patterns Change protected paths from `/api/products/create` to `/api/products` with method checks, or handle auth inside route handlers consistently (pick one pattern).

1c. Lock CORS to production domain Replace `'Access-Control-Allow-Origin': '*'` with `process.env.ALLOWED_ORIGIN`.

1d. Create `Dockerfile.worker` Worker container won't build without it — BullMQ jobs won't run.

1e. Add Docker healthchecks Prevent app from starting before Postgres/Qdrant/Redis are ready.

Priority 2 — HIGH (Within 2 Weeks of Launch)

2a. Cache embedding results in Redis

```
key: "embed:" + sha256(query)  
ttl: 3600 seconds
```

Eliminates ~300ms latency on repeated/similar queries.

2b. Cache preprocessed queries

```
key: "preprocess:" + sha256(rawQuery + language)  
ttl: 3600 seconds
```

Eliminates one full GPT-4o call per search.

2c. Add score threshold to Qdrant search Add `score_threshold: 0.3` to filter out low-confidence results — prevents showing unrelated products.

2d. Add Prisma connection pool config

```
DATABASE_URL="...?connection_limit=20&pool_timeout=30"
```

2e. Move S3 to AWS Bahrain (me-south-1) Reduces image load latency from ~250ms to ~30ms for Kuwait users.

2f. Add file upload validation Check MIME type and enforce max size (e.g. 10MB) server-side in the presign endpoint.



Priority 3 — MEDIUM (Month 1 Post-Launch)

3a. Add cross-encoder reranking Integrate `Cohere Rerank` or a lightweight `BGE-reranker` for top-50 → top-10 reranking. This dramatically improves result quality for ambiguous queries.

3b. Add Arabic stemming to BM25 tokenizer Replace simple tokenization with `arabic-stemmer` library — handles morphological variants like سيارة/سيارات.

3c. Implement PM2 cluster mode

```
// ecosystem.config.js
instances: 'max', // use all CPU cores
exec_mode: 'cluster'
```

3d. Add score_threshold to Qdrant + "no results" graceful handling Currently returns 0-count products silently — show a clear "No listings found, be the first to post one!" CTA.

3e. Complete localization (Sprint 6 remaining 40%) Error messages, validation strings, edge case text — needed before Arabic-speaking users see English error messages.

3f. Database connection via PgBouncer For 100k+ product scale, add connection pooler between Next.js and Postgres.

Priority 4 — LOW (Month 2-3)

4a. Product expiry job `expiresAt` field exists (`NOW() + 30 days`) but no cron/BullMQ job to mark expired products. Run daily.

4b. Add CloudFront CDN for S3 images Cache images at edge — reduces S3 egress cost and improves load time globally.

4c. Qdrant snapshot backup Add scheduled Qdrant snapshot API calls (daily) — losing the vector index means re-embedding all products.

4d. Monitoring stack Add Grafana + Prometheus or use a SaaS (Datadog, New Relic) — track search latency, error rates, embedding failures.

4e. Add favorites/watchlist Post-V1 feature — `Favorite` model is missing from schema but the UX has placeholders.

10. FINAL ASSESSMENT SUMMARY

COMPONENT GRADES	
AI Search Pipeline	A- (excellent)
Database Schema & Design	A (excellent)
Auth & JWT Implementation	B+ (solid)



API Route Coverage	A- (complete)
Frontend / Chat UX	B+ (complete)
SSE Streaming Implementation	A (excellent)
Hybrid Search (Dense + BM25 + RRF)	A (excellent)
Security Posture	C+ (gaps)
Deployment Readiness	C (not ready)
Test Coverage	D (missing)
Localization	B- (partial)
Scalability Architecture	B (good base)
Code Quality & Consistency	B+ (solid)

OVERALL: B+ — Strong core with specific launch blockers

The system is architecturally excellent — the hybrid search design (dense + BM25 + RRF + preprocessing + reranking) is far more sophisticated than typical MVP-level work. The main risks are in the deployment layer (missing files, dimension bug, CORS) and missing test coverage. Fix the 5 Priority 1 items and this system is ready for a guarded launch.

The background agent hit a permissions issue, but it doesn't matter — I read all those files directly in the main conversation and the full audit report above is already based on the actual code.

Let me save the key findings to memory for future sessions: